

# Homework 08

## ⚠ Before you start ⚠





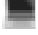

*Duplicate this Jupyter Notebook in your `week-10` folder (right-click -> Duplicate) and then add your last name to the beginning of it (ie. `blevins-hw-08.ipynb` - otherwise you risk having all your work overwritten when you try to sync your GitHub repository with your instructor's repository.*

---

## Overview

In this assignment, you'll synthesize some of the Python skills you've learned over the past month or so, including Pandas and Plotly. You'll be analyzing the opening of new businesses in Colorado during the 1940s.

Draw on the following tutorials:

-  Walsh, [Pandas Basics Part 1](#)
-  Walsh, [Pandas Basics Part 2](#)
-  Walsh, [Pandas Basics Part 3](#)
-  [Pandas Concepts](#)
-  [Introduction to Plotly](#)
-  [Cleaning Excel Files](#)

## The Data

First, get the necessary data files from our shared course repository:

- Open GitHub Desktop and select your course repository ( `lastname-sp25-data-materials` )
- Click `Fetch origin` to check for updates
- Go to `Branch` → `Merge into current branch` → select `upstream/main` -> `Merge`
- Click `Push origin` to sync everything up
- Launch Jupyter Lab and navigate to the `week-10` folder

You should see a single Excel file that you will be working with: `co-new-businesses-1940s.xlsx` . Inside that Excel file, there are two separate sheets: `New CO Businesses` and `Cities 1940` .

- **New CO Businesses** : This is a subset of new businesses that were established in Colorado during the 1940s - a subset of data drawn from [this database](#).
- **Cities 1940** : this contains population statistics for Colorado cities in the 1940 Census.

## Import Libraries and Load Data

- Import the necessary libraries:
  - pandas (using the alias `pd` )
  - plotly.express (using the alias `px` )

```
In [7]: import pandas as pd
import plotly.express as px
```

- Load both sheets from the Excel file:
  - Create a variable called `businesses_df` to store the "New CO Businesses" sheet in the Excel file
  - Create a variable called `cities_df` to store the "Cities 1940" sheet in the Excel file
  - Use `pd.read_excel()` with the appropriate parameters

```
In [9]: businesses_df= pd.read_excel("co-new-businesses-1940s.xlsx", sheet_name= "Ne
cities_df = pd.read_excel("co-new-businesses-1940s.xlsx", sheet_name= 'Citie
```

## Familiarize Yourself with the Data

Familiarize yourself with the data:

- Display a sample of 10 rows from each dataframe.
- Check the data types for the columns in each dataframe

```
In [11]: print(f"{businesses_df.sample(10)}\n")

print(f"{cities_df.sample(10)}\n")

print(f"{cities_df.dtypes}\n")

print(f"{businesses_df.dtypes}\n")
```

	entityid	Business entity name \
722	19871116551	MIDDLE PARK WATER USERS PROTECTIVE ASSOCIATION
952	19871110200	CHURCH OF THE ASCENSION & HOLY TRINITY & ASCEN...
4	19871117433	LYNCH-COTTEN POST NO. 190, THE AMERICAN LEGION
834	19871313658	PIONEER ASTRO INDUSTRIES, INC., Dissolved July...
334	19491116218	Morgan Game and Fish Conservation Club
767	19871141764	RED ROCKS BAPTIST CHURCH
354	19871007393	MELVILLE REALTY COMPANY, INC., Colorado Author...
624	19871113775	UNITED LUMBER AND HARDWARE, INC., Delinquent J...
484	19871109816	THE SOUTH PARK POST 172-THE AMERICAN LEGION
578	19871112651	KIRK OF BONNIE BRAE (CONGREGATIONAL)

	Address	city	state	zip_code	Country \
722	NaN	NaN	NaN	NaN	NaN
952	420 W 18th St	Pueblo	CO	81003.0	US
4	425 Highway 92	Crawford	CO	81415.0	US
834	3410 N PROSPECT ST	COLORADO SPRINGS	CO	80907.0	US
334	427 S. Lake Street	Fort Morgan	CO	80701.0	US
767	14711 W Morrison Rd	Morrison	CO	80465.0	US
354	1 CVS DR	WOONSOCKET	RI	2895.0	US
624	307 E BRIDGE ST	BRIGHTON	CO	80601.0	US
484	602 Clark St	FAIRPLAY	CO	80440.0	US
578	1201 S Steele St	Denver	CO	80210.0	US

	date_entity_formed	year_entity_formed
722	1949-06-06	1949
952	1946-01-30	1946
4	1949-12-30	1949
834	1945-01-08	1945
334	1949-04-11	1949
767	1945-01-11	1945
354	1945-07-12	1945
624	1947-12-30	1947
484	1945-10-24	1945
578	1947-05-08	1947

	city	year	total population
47	delta	1940	3717
185	sedgwick	1940	373
202	timnath	1940	147
182	saguache	1940	1219
119	lamar	1940	4445
88	grover	1940	137
106	jamestown	1940	190
132	manitou springs	1940	1462
81	glenwood springs	1940	2253
73	fowler	1940	922

city	object
year	int64
total population	int64

dtype: object

entityid	int64
Business entity name	object
Address	object

```

city                object
state               object
zip_code            float64
Country             object
date_entity_formed  object
year_entity_formed  int64
dtype: object

```

## Data Cleaning and Preparation

### Cleaning column names

For both datasets, you want to clean and standardize the column names (headers):

- Change column names to all lowercase
- Replace any whitespace with an underscore ( `_` ) - ex. `some column` becomes `some_column`
- *Hint: Use `str.lower()` and `str.replace()`*
- Show the first 10 rows of your dataframe to make sure it worked

```

In [13]: #Your code here
businesses_df.columns = [col.strip().lower().replace(' ', '_').replace('-', '_') for col in businesses_df.columns]
cities_df.columns = [col.strip().lower().replace(' ', '_').replace('-', '_') for col in cities_df.columns]

businesses_df.head(10)
cities_df.head(10)

```

```

Out[13]:
   city  year  total_population
0  akron  1940                1417
1  alamosa  1940               5613
2    alma  1940                469
3  antonito  1940               1220
4   arriba  1940                286
5  arvada  1940               1482
6   aspen  1940                777
7  aurora  1940               3437
8   basalt  1940                212
9  bayfield  1940                372

```

### Standardize and clean data for cities

- Standardize city names in the business data so that it **removes any trailing or**

**leading whitespace** and **changes the values to all lowercase** (hint: use

`.str.strip()` and `.str.lower()` )

- Show the first 10 rows of your dataframe to make sure it worked

In [15]: *#Your code here*

```
businesses_df["city"] = businesses_df['city'].str.strip().str.lower()

businesses_df.head(10)
```

Out[15]:

	entityid	business_entity_name	address	city	state	zip_code	country
0	19871004753	ALAMOSA CREDIT UNION	2437 MAIN ST	alamosa	CO	81101.0	US
1	19871241137	THE UNITED METHODIST CHURCH OF STEAMBOAT SPRINGS	736 OAK ST	steamboat springs	CO	80487.0	US
2	19871275274	ALLIED JEWISH FEDERATION OF COLORADO	300 S. Dahlia St.	denver	CO	80246.0	US
3	19871127721	Iglesia CRISTO REY + Christ the King, ELCA	2300 S Patton Ct	denver	CO	80219.0	US
4	19871117433	LYNCH-COTTEN POST NO. 190, THE AMERICAN LEGION	425 Highway 92	crawford	CO	81415.0	US
5	19871105155	THE BEAR RIVER VALLEY FARMERS COOPERATIVE	193 E Jefferson Ave	hayden	CO	81639.0	US
6	19871162072	Belmar Baptist Church	460 S Kipling St	lakewood	CO	80226.0	US
7	19871110810	Bethel Lutheran Church of Windsor, Colorado	328 Walnut St	windsor	CO	80550.0	US
8	19871116977	BLACKINTON AND DECKER, INC., Delinquent Novemb...	424 LIPAN	denver	CO	80204.0	US
9	19871113871	BOW-MAR OWNERS, INC.	5380 Lakeshore Dr	littleton	CO	80123.0	US

## Categorize Cities

### Define your function

Create a function called `categorize_city_size` that does the following:

- Takes in a number that corresponds to the population for a city and returns the following based on the size of the city:
  - Small Town if population is less than 1,000
  - Medium Town if population is between 1,000 to 5,000
  - Large Town if population is between 5,000 to 20,000
  - City if population greater than or equal to 20,000

```
In [17]: def categorize_city_size(population):  
        if population >= 20000:  
            return "City"  
        elif population > 5000 and population < 20000:  
            return "Large Town"  
        elif population > 1000 and population < 5000:  
            return "Medium Town"  
        else:  
            return "Small Town"
```

## Test Your Function

Test out the function on a single number ( 2,000 ) to make sure it returns Medium Town

```
In [19]: categorize_city_size(2000)  
  
cities_df.head(10)
```

```
Out[19]:
```

	city	year	total_population
0	akron	1940	1417
1	alamosa	1940	5613
2	alma	1940	469
3	antonito	1940	1220
4	arriba	1940	286
5	arvada	1940	1482
6	aspen	1940	777
7	aurora	1940	3437
8	basalt	1940	212
9	bayfield	1940	372

## Apply the function

- Take your `cities_df` dataframe and add a new column called `city_category` that applies your function to the `total_population` column of the dataframe.
- *Hint: use `apply()`*
- Show the first 10 rows of your dataframe to make sure it worked

```
In [21]: cities_df['city_category'] = cities_df['total_population'].apply(categorize_cities_df.head(10))

businesses_df.head(10)
```

```
Out[21]:
```

	entityid	business_entity_name	address	city	state	zip_code	country
0	19871004753	ALAMOSA CREDIT UNION	2437 MAIN ST	alamosa	CO	81101.0	US
1	19871241137	THE UNITED METHODIST CHURCH OF STEAMBOAT SPRINGS	736 OAK ST	steamboat springs	CO	80487.0	US
2	19871275274	ALLIED JEWISH FEDERATION OF COLORADO	300 S. Dahlia St.	denver	CO	80246.0	US
3	19871127721	Iglesia CRISTO REY + Christ the King, ELCA	2300 S Patton Ct	denver	CO	80219.0	US
4	19871117433	LYNCH-COTTEN POST NO. 190, THE AMERICAN LEGION	425 Highway 92	crawford	CO	81415.0	US
5	19871105155	THE BEAR RIVER VALLEY FARMERS COOPERATIVE	193 E Jefferson Ave	hayden	CO	81639.0	US
6	19871162072	Belmar Baptist Church	460 S Kipling St	lakewood	CO	80226.0	US
7	19871110810	Bethel Lutheran Church of Windsor, Colorado	328 Walnut St	windsor	CO	80550.0	US
8	19871116977	BLACKINTON AND DECKER, INC., Delinquent Novemb...	424 LIPAN	denver	CO	80204.0	US
9	19871113871	BOW-MAR OWNERS, INC.	5380 Lakeshore Dr	littleton	CO	80123.0	US

## Analyze Businesses by Year

Let's take a look at how many new businesses were formed in Colorado in each year during the 1940s:

## Calculate new businesses by year

Create a variable called `businesses_per_year` by:

- Counting the number of new businesses based on `year_entity_formed`
- *Hint: use `value_counts()` and `reset_index()`*
- Show the first 10 rows of your dataframe

```
In [23]: businesses_per_year= businesses_df[['year_entity_formed']].value_counts().re  
businesses_per_year.head(10)
```

```
Out[23]:
```

	year_entity_formed	count
0	1947	161
1	1948	156
2	1946	153
3	1949	133
4	1945	87
5	1940	72
6	1941	69
7	1943	47
8	1944	43
9	1942	35

## Visualize new businesses by year

Create a bar chart using Plotly Express showing new businesses per year:

- Set x-axis to the year
- Set y-axis to the number of new businesses
- Add an appropriate title and labels
- Display text on each bar
- Hint: Use `px.bar()`

```
In [25]: fig = px.bar(  
    businesses_per_year,  
    x='year_entity_formed',  
    y='count',  
    title='Number of Businesses Formed in Colorado Per Year', # Add a title  
    labels={'year_entity_formed': 'year', 'count': '# of Businesses'}, # Rena  
    color='count',  
    color_continuous_scale='Agsunset_r',
```



```
        template='xgridoff'  
    )  
fig.show()
```

## Analyze Businesses by City

Let's take a look at how many new businesses were formed in each Colorado city during the 1940s:

## Calculate number of new businesses by city

Create a new variable called `city_businesses` that contains:

- A dataframe with counts of the number of new businesses in each city
- *Hint: Use `value_counts()` and `reset_index()`*
- Show the first 10 rows of your dataframe

In [27]: *#Your code here*

```
city_businesses= businesses_df["city"].value_counts().reset_index()
city_businesses.head(10)
```

Out[27]:

	city	count
0	denver	152
1	colorado springs	34
2	lakewood	22
3	pueblo	20
4	arvada	14
5	grand junction	14
6	fort collins	13
7	greeley	13
8	centennial	12
9	englewood	12

## Visualize new businesses by city

Create a bar chart with Plotly Express showing the top 10 cities with the most new businesses created during the 1940s:

- Filter to only show the top 10 cities (hint: use `.head()` )
- Set x-axis to `city`
- Set y-axis to `count`
- Add an appropriate title and labels

```
In [29]: fig = px.bar(
    city_businesses.head(10),
    x='city',
    y='count',
    title='Top 10 Cities With the Most New Businesses during the 1940s', #
    labels={'city': 'City', 'count': '# of Businesses'}, # Rename axis labels
    color='city',
    color_continuous_scale='Agsunset_r',
    template='xgridoff'

)


fig.show()
```

## Combine Business and City Data

We have two datasets, both of which contain information about Colorado cities. Let's combine the two into a single dataframe that contains both information about new businesses and their population in the 1940 census.

### Merge dataframes

Merge the two dataframes together:

- Create a new variable called `merged_df`
- Use `pd.merge()` on the `city_businesses` and `cities_df` dataframes
- Figure out which column is shared between the two to use as your "key" to merge them
-  **Note: use the `how='inner'` parameter for your merge**
- Show the first 10 rows of your new dataframe

```
In [31]: #Your code here
```

```
merged_df = pd.merge(
    city_businesses,
    cities_df,
    on='city',
    how='inner'
)

merged_df.head(10)
```

Out[31]:

	city	count	year	total_population	city_category
0	denver	152	1940	322412	City
1	colorado springs	34	1940	36789	City
2	pueblo	20	1940	52162	City
3	arvada	14	1940	1482	Medium Town
4	grand junction	14	1940	12479	Large Town
5	fort collins	13	1940	12251	Large Town
6	greeley	13	1940	15995	Large Town
7	englewood	12	1940	9680	Large Town
8	littleton	11	1940	2244	Medium Town
9	aurora	10	1940	3437	Medium Town

## Filter out missing values

You'll note that several rows of data contain `NaN` or missing values - this means that there was a city listed in the businesses dataframe but it didn't have a corresponding match in the population dataframe. For now, remove these from the `merged_df` dataframe:

- Filter out rows where `total_population` is `NaN`
- *Hint: use a filter + `.notna()`*

```
In [33]: merged_df = merged_df[merged_df['total_population'].notna()]
```

## Calculate new businesses on a per capita rate

To make it easier to compare larger cities with smaller cities, you're going to calculate a new column for each city: the number of new businesses per 1,000 residents.

- Add a new column to `merged_df` called `biz_per_thousand` that is filled with:
  - A calculation dividing the `count` column by the `total_population` column and multiplying by 1,000
- Sort the merged dataframe by `biz_per_thousand` in descending order

- Show the first 10 rows of the dataframe to check if it worked

```
In [35]: merged_df['biz_per_thousand'] = merged_df['count']/merged_df['total_population']
merged_df = merged_df.sort_values(by='biz_per_thousand', ascending=False)
merged_df.head(10)
```

```
Out[35]:
```

	city	count	year	total_population	city_category	biz_per_thousand
92	green mountain falls	1	1940	87	Small Town	11.494253
36	keenesburg	3	1940	284	Small Town	10.563380
59	bennett	2	1940	199	Small Town	10.050251
3	arvada	14	1940	1482	Medium Town	9.446694
52	morrison	2	1940	216	Small Town	9.259259
20	castle rock	5	1940	580	Small Town	8.620690
33	woodland park	3	1940	372	Small Town	8.064516
54	granby	2	1940	251	Small Town	7.968127
79	grover	1	1940	137	Small Town	7.299270
96	timnath	1	1940	147	Small Town	6.802721

## Visualize new business creation by city

Let's say we want to see the cities with the highest *rate* of business creation (ie. new businesses per thousand residents)

- Create a bar chart in Plotly of `merged_df` :
  - Filter to only show the top 10 cities (use `.head(10)` )
  - Set x-axis to `city`
  - Set y-axis to `biz_per_thousand`
  - Use `city_category` for color
  - Add an appropriate title and labels

```
In [37]: fig = px.bar(
merged_df.head(10),
x='city',
y='biz_per_thousand',
title='Cities with highest new businesses per thousand residents', # Add title
labels={'city': 'City', 'biz_per_thousand': '# of Businesses per 1000 residents'},
color='city_category',
color_continuous_scale='Inferno',
template='xgridoff')
```

```
)  
  
fig.show()
```

## Bonus: New businesses by city category

Let's say we want to compare different size categories to see whether new businesses were cropping up in smaller places or bigger cities.

### Create a new dataframe

First, you'll need to create a new dataframe that consists of four rows, with each row a different category of city containing the total number of businesses created within that category of city.

- Create a new dataframe called `city_category_totals`
- Start with `merged_df`
- Group by `city_category`
- Add up ( `sum()` ) the `count` column

- Use `.reset_index()`

In [39]: `#Your code here`

## Visualize businesses by city category

- Create a [pie chart](#) in Plotly:
  - Use `px.pie()` with appropriate parameters
  - Use `city_category_totals` as your dataframe
  - Use `count` for your values
  - Use `city_category` for your names
  - Add an appropriate title and labels

In [41]: `#Your code here`

## Bonus Challenge: Create a Scatterplot

Create a scatter plot in Plotly showing:

- The relationship between city population (x-axis) and new businesses (y-axis)
- Only data for towns with a population of 2,000 or more people.
- Dots sized according to the number of new businesses in that city
- Dots colored according to their size category

In [43]: `#Your code here`

## Submission Guidelines

- Run all code cells and make sure it is outputting without errors
- Submit both the notebook file (.ipynb) and a PDF export of your notebook [on Canvas](#)
- Note: the PDF probably won't display the Plotly figures - that's okay